

1)

The computer I use runs macOS Mojave 10.14.3 as its Operating System. Hence, I had to choose between allocating a portion of my hard-disk space to Ubuntu 18.04.1 LTS or to use it via a Virtual Machine. I chose the latter option and installed Oracle's VirtualBox 6.0.4. VirtualBox is essentially an efficient x86 and AMD64/Intel64 virtualization product. It is developed and regularly maintained by Oracle, so I found it reliable.

The installation process was very time-consuming and troublesome for me. I went through some video tutorials on the installation process; however, none of the recorded solutions worked in my computer. Thus, I tried the exact same steps in another macOS Mojave computer. Surprisingly, it worked.

The 10 Linux commands I learned from the Linux usage tutorials are as follows:

- `ls`: Shows all the major directories under a specified file system.
- `cd`: Lets the user to change between directories.
- `mv`: Lets the user to move a specified file to a particular directory.
- `mkdir`: Creates a new directory.
- `rmdir`: Removes an existing directory.
- `touch`: Creates a new (empty) file.
- `rm`: Removes an existing file.
- `clear`: Clears the screen and takes the user back to the initial prompt of the working directory.
- `locate`: Helps finding a particular file.
- `history`: Shows all commands used in the current session.

2)

(i) The kernel executable is named `vmlinuz` and is located under: `/boot`

(ii) Output of `uname -r` command (Version number): `4.15.0-29-generic`

3)

I downloaded the `4.14.99` version from `kernel.org`'s archive, since it was the closest numbered version.

The subdirectories of the root directory of the source code (`/linux-4.14.99/`) are as follows:

- arch
- block
- certs
- crypto
- Documentation
- drivers
- firmware
- fs
- include
- init
- ipc
- kernel
- lib
- mm
- net
- samples
- scripts
- security
- sound
- tools
- usr
- virt

4)

The system call table definition in the linux-4.14.99 source code has the path:

`/linux-4.14.99/arch/x86/entry/syscalls/syscall_64.tbl`

System call names corresponding to the asked system call numbers are given below:

- 5: `fstat`
- 43: `accept`
- 123: `setfsgid`

- 220: semtimedop

5)

strace records the names of the system calls that are called by the specified command until the command exits, together with the arguments and return values of the system calls.

Sample output for **strace ls** is provided, starting from below:

```

1 efe@efe-VirtualBox:~$ strace ls
2 execve("/bin/ls", ["ls"], 0x7ffddbfb811b0 /* 61 vars */) = 0
3 brk(NULL) = 0x55e0cae5d000
4 access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
5 access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
6 openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
7 fstat(3, {st_mode=S_IFREG|0644, st_size=73553, ...}) = 0
8 mmap(NULL, 73553, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7ff3f7088000
9 close(3) = 0
10 access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
11 openat(AT_FDCWD, "/lib/x86_64-linux-gnu/libselinux.so.1", O_RDONLY|O_CLOEXEC) = 3
12 read(3, "\177ELF\2\1\1\0\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0\20b\0\0\0\0\0"... , 832)
    = 832
13 fstat(3, {st_mode=S_IFREG|0644, st_size=154832, ...}) = 0
14 mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0
    x7ff3f7086000
15 mmap(NULL, 2259152, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0
    x7ff3f6c4b000
16 mprotect(0x7ff3f6c70000, 2093056, PROT_NONE) = 0
17 mmap(0x7ff3f6e6f000, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|
    MAP_DENYWRITE, 3, 0x24000) = 0x7ff3f6e6f000
18 mmap(0x7ff3f6e71000, 6352, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|
    MAP_ANONYMOUS, -1, 0) = 0x7ff3f6e71000
19 close(3) = 0
20 access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
21 openat(AT_FDCWD, "/lib/x86_64-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
22 read(3, "\177ELF\2\1\1\3\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0\260\34\2\0\0\0\0"... ,
    832) = 832
23 fstat(3, {st_mode=S_IFREG|0755, st_size=2030544, ...}) = 0
24 mmap(NULL, 4131552, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0
    x7ff3f685a000
25 mprotect(0x7ff3f6a41000, 2097152, PROT_NONE) = 0
26 mmap(0x7ff3f6c41000, 24576, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|
    MAP_DENYWRITE, 3, 0x1e7000) = 0x7ff3f6c41000
27 mmap(0x7ff3f6c47000, 15072, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|
    MAP_ANONYMOUS, -1, 0) = 0x7ff3f6c47000

```

```
28 close(3) = 0
29 access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
30 openat(AT_FDCWD, "/lib/x86_64-linux-gnu/libpcres.so.3", O_RDONLY|O_CLOEXEC) = 3
31 read(3, "\177ELF\2\1\1\0\0\0\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0 \25\0\0\0\0\0\0"... , 832)
   = 832
32 fstat(3, {st_mode=S_IFREG|0644, st_size=464824, ...}) = 0
33 mmap(NULL, 2560264, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0
   x7ff3f65e8000
34 mprotect(0x7ff3f65e8000, 2097152, PROT_NONE) = 0
35 mmap(0x7ff3f6858000, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|
   MAP_DENYWRITE, 3, 0x70000) = 0x7ff3f6858000
36 close(3) = 0
37 access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
38 openat(AT_FDCWD, "/lib/x86_64-linux-gnu/libdl.so.2", O_RDONLY|O_CLOEXEC) = 3
39 read(3, "\177ELF\2\1\1\0\0\0\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0P\16\0\0\0\0\0\0"... , 832)
   = 832
40 fstat(3, {st_mode=S_IFREG|0644, st_size=14560, ...}) = 0
41 mmap(NULL, 2109712, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0
   x7ff3f63e4000
42 mprotect(0x7ff3f63e7000, 2093056, PROT_NONE) = 0
43 mmap(0x7ff3f65e6000, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|
   MAP_DENYWRITE, 3, 0x2000) = 0x7ff3f65e6000
44 close(3) = 0
45 access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
46 openat(AT_FDCWD, "/lib/x86_64-linux-gnu/libpthread.so.0", O_RDONLY|O_CLOEXEC) = 3
47 read(3, "\177ELF\2\1\1\0\0\0\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0000b\0\0\0\0\0\0"... , 832)
   = 832
48 fstat(3, {st_mode=S_IFREG|0755, st_size=144976, ...}) = 0
49 mmap(NULL, 2221184, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0
   x7ff3f61c5000
50 mprotect(0x7ff3f61df000, 2093056, PROT_NONE) = 0
51 mmap(0x7ff3f63de000, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|
   MAP_DENYWRITE, 3, 0x19000) = 0x7ff3f63de000
52 mmap(0x7ff3f63e0000, 13440, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|
   MAP_ANONYMOUS, -1, 0) = 0x7ff3f63e0000
53 close(3) = 0
54 mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0
   x7ff3f7084000
55 arch_prctl(ARCH_SET_FS, 0x7ff3f7085040) = 0
56 mprotect(0x7ff3f6c41000, 16384, PROT_READ) = 0
57 mprotect(0x7ff3f63de000, 4096, PROT_READ) = 0
58 mprotect(0x7ff3f65e6000, 4096, PROT_READ) = 0
59 mprotect(0x7ff3f6858000, 4096, PROT_READ) = 0
60 mprotect(0x7ff3f6e6f000, 4096, PROT_READ) = 0
```

```

61 mprotect(0x55e0c98d5000, 8192, PROT_READ) = 0
62 mprotect(0x7ff3f709a000, 4096, PROT_READ) = 0
63 munmap(0x7ff3f7088000, 73553) = 0
64 set_tid_address(0x7ff3f7085310) = 2097
65 set_robust_list(0x7ff3f7085320, 24) = 0
66 rt_sigaction(SIGRTMIN, {sa_handler=0x7ff3f61cacb0, sa_mask=[], sa_flags=
    SA_RESTORER|SA_SIGINFO, sa_restorer=0x7ff3f61d7890}, NULL, 8) = 0
67 rt_sigaction(SIGRT_1, {sa_handler=0x7ff3f61cad50, sa_mask=[], sa_flags=SA_RESTORER
    |SA_RESTART|SA_SIGINFO, sa_restorer=0x7ff3f61d7890}, NULL, 8) = 0
68 rt_sigprocmask(SIG_UNBLOCK, [RTMIN RT_1], NULL, 8) = 0
69 prlimit64(0, RLIMIT_STACK, NULL, {rlim_cur=8192*1024, rlim_max=RLIM64_INFINITY}) =
    0
70 statfs("/sys/fs/selinux", 0x7ffd3bd51d50) = -1 ENOENT (No such file or directory)
71 statfs("/selinux", 0x7ffd3bd51d50) = -1 ENOENT (No such file or directory)
72 brk(NULL) = 0x55e0cae5d000
73 brk(0x55e0cae7e000) = 0x55e0cae7e000
74 openat(AT_FDCWD, "/proc/filesystems", O_RDONLY|O_CLOEXEC) = 3
75 fstat(3, {st_mode=S_IFREG|0444, st_size=0, ...}) = 0
76 read(3, "nodev\tsysfs\nnodev\trootfs\nnodev\ttr"... , 1024) = 383
77 read(3, "", 1024) = 0
78 close(3) = 0
79 access("/etc/selinux/config", F_OK) = -1 ENOENT (No such file or directory)
80 openat(AT_FDCWD, "/usr/lib/locale/locale-archive", O_RDONLY|O_CLOEXEC) = 3
81 fstat(3, {st_mode=S_IFREG|0644, st_size=11731760, ...}) = 0
82 mmap(NULL, 11731760, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7ff3f5694000
83 close(3) = 0
84 ioctl(1, TCGETS, {B38400 opost isig icanon echo ...}) = 0
85 ioctl(1, TIOCGWINSZ, {ws_row=28, ws_col=79, ws_xpixel=0, ws_ypixel=0}) = 0
86 openat(AT_FDCWD, ".", O_RDONLY|O_NONBLOCK|O_CLOEXEC|O_DIRECTORY) = 3
87 fstat(3, {st_mode=S_IFDIR|0755, st_size=4096, ...}) = 0
88 getdents(3, /* 21 entries */, 32768) = 664
89 getdents(3, /* 0 entries */, 32768) = 0
90 close(3) = 0
91 fstat(1, {st_mode=S_IFCHR|0620, st_rdev=makedev(136, 0), ...}) = 0
92 write(1, "Desktop Downloads\t Music "..., 52Desktop Downloads Music Public Videos
93 ) = 52
94 write(1, "Documents examples.desktop Pic"... , 49Documents examples.desktop
    Pictures Templates
95 ) = 49
96 close(1) = 0
97 close(2) = 0
98 exit_group(0) = ?
99 +++ exited with 0 +++

```

6)

`time` command provides timing statistics about the execution of a specified command, the different times the command outputs have the following meanings:

- **real**: The total time from start to finish of the call, ie. from the moment you press Enter until the moment the command returns.
- **user**: The amount of CPU time that is spent in the user mode, in which privileged instructions cannot be executed.
- **sys**: The amount of CPU time that is spent in the kernel mode, in which privileged instructions are allowed to be executed.

The following table displays the timing statistics for five different commands:

command (including arguments)	real	user	sys
cp Music	0.016s	0.002s	0.000s
time strace ls	0.007s	0.002s	0.003s
cd Desktop	0.000s	0.000s	0.000s
ls	0.003s	0.001s	0.000s
time time strace ls	0.006s	0.005s	0.000s

Table 1: Some commands and their timing statistics

7)

I measured the execution times for `getpid()`, `open()`, `read()`, `write()` and `mkdir()` system calls. `getpid()` has a simple job and it does not have any arguments, so it had a relatively short execution time as expected. I used `open()` system call to create a non-existing `.txt` file with the permissions necessary to perform read and write operations. The execution for this took relatively long, probably because the kernel had to allocate many resources under the hood to create the file. Then I measured the execution times for `write()` and `read()` system calls for various number of bytes. The results made sense, since the execution time increased with the number of bytes. The very first execution of `write()` took longer than I have expected, this may again be something about the kernel internals. I intentionally called `write()`s before `read()`s to fill the `.txt` file I had created with enough number of bytes. An implementation detail was to close and re-open the `.txt` file in between the write and read operations, otherwise the pointer used to read the file gets stuck in the end-of-file byte. Lastly, I made time measurements on `mkdir()` system call for two cases. In one case, a directory was created with full permissions; and in the other, a directory was created with only read permissions. The case with full permissions had a longer execution time as expected, since the kernel performs specific operations to grant those permissions.

The experiment is summarized in the table below:

system call (arguments)	execution time (in microseconds)
getpid()	2
open("read.txt", O_CREAT O_RDWR O_APPEND, 00700)	33
write(fileDescriptor, &bytes, 100)	13
write(fileDescriptor, &bytes, 1000)	3
write(fileDescriptor, &bytes, 10000)	7
write(fileDescriptor, &bytes, 100000)	42
read(fileDescriptor, &buffer, 100)	1
read(fileDescriptor, &buffer, 1000)	1
read(fileDescriptor, &buffer, 10000)	4
read(fileDescriptor, &buffer, 100000)	31
mkdir("allAccessPermissionsGiven", ACCESSPERMS)	25
mkdir("onlyReadPermissionsGiven", S_IRUSR S_IRGRP S_IROTH)	14

Table 2: Some system calls and their execution times

The C code and command line output follows:

The C code:

```

100 /**
101  * This program performs some timing experiments on some of the commonly
102  * used Linux system calls. It executes the system calls with the specified
103  * arguments and reports the execution time in microseconds.
104  * @author Efe Acer
105  * @version 1.0
106  */
107
108 // Necessary imports to be able to run the system calls
109 #include <sys/time.h>
110 #include <sys/types.h>
111 #include <unistd.h>
112 #include <fcntl.h>
113 #include <sys/stat.h>
114 #include <stdio.h>
115
116 // Function declaration(s)
117 unsigned long getCurrentTime();
118
119 int main() {
120     printf("\nTime measurements for different system calls in microseconds follows

```

```
        :\n");
121
122 // Preparation for getpid() measurements
123 unsigned long getpidStart;
124 unsigned long getpidEnd;
125 int processID;
126
127 // Get process ID
128 getpidStart = getCurrentTime();
129 processID = getpid();
130 getpidEnd = getCurrentTime();
131 printf("\nTime to execute getpid(): %ld\n", getpidEnd - getpidStart);
132 printf("Process ID is: %d\n", processID);
133
134 // Preparation for open() measurements
135 int fileDescriptor;
136 unsigned long openStart;
137 unsigned long openEnd;
138
139 // Create a .txt file that will be used to read bytes from
140 openStart = getCurrentTime();
141 fileDescriptor = open("read.txt", O_CREAT | O_RDWR | O_APPEND, 00700); //
    00700 is for file owner permissions
142 openEnd = getCurrentTime();
143 printf("\nTime to execute open() to create a new .txt file is: %ld\n", openEnd
    - openStart);
144 printf("The value of fileDescriptor is %d:\n", fileDescriptor);
145
146 // Preparation for write() measurements
147 unsigned char bytes[100000];
148 for (int i = 0; i < 100000; i++) {
149     bytes[i] = '.';
150 }
151 unsigned long writeStart;
152 unsigned long writeEnd;
153 int numWritten;
154
155 // Write 100 bytes to read.txt
156 writeStart = getCurrentTime();
157 numWritten = write(fileDescriptor, &bytes, 100);
158 writeEnd = getCurrentTime();
159 printf("\nTime to execute write() for %d bytes: %ld\n", numWritten, (writeEnd
    - writeStart));
160
```



```
161 // Write 1000 bytes to read.txt
162 writeStart = getCurrentTime();
163 numWritten = write(fileDescriptor, &bytes, 1000);
164 writeEnd = getCurrentTime();
165 printf("Time to execute write() for %d bytes: %ld\n", numWritten, (writeEnd -
    writeStart));
166
167 // Write 10000 bytes to read.txt
168 writeStart = getCurrentTime();
169 numWritten = write(fileDescriptor, &bytes, 10000);
170 writeEnd = getCurrentTime();
171 printf("Time to execute write() for %d bytes: %ld\n", numWritten, (writeEnd -
    writeStart));
172
173 // Write 100000 bytes to read.txt
174 writeStart = getCurrentTime();
175 numWritten = write(fileDescriptor, &bytes, 100000);
176 writeEnd = getCurrentTime();
177 printf("Time to execute write() for %d bytes: %ld\n", numWritten, (writeEnd -
    writeStart));
178
179 // Restore the pointer that is used to read and write to the file
180 close(fileDescriptor);
181 fileDescriptor = open("read.txt", O_RDWR, 00700); // 00700 is for file owner
    permissions
182
183 // Preparation for read() measurements
184 unsigned char buffer[111100];
185 unsigned long readStart;
186 unsigned long readEnd;
187 int numRead;
188
189 // Read 100 bytes from read.txt
190 readStart = getCurrentTime();
191 numRead = read(fileDescriptor, &buffer, 100);
192 readEnd = getCurrentTime();
193 printf("\nTime to execute read() for %d bytes: %ld\n", numRead, (readEnd -
    readStart));
194
195 // Read 1000 bytes from read.txt
196 readStart = getCurrentTime();
197 numRead = read(fileDescriptor, &buffer, 1000);
198 readEnd = getCurrentTime();
199 printf("Time to execute read() for %d bytes: %ld\n", numRead, (readEnd -
```

```
        readStart));
200
201    // Read 10000 bytes from read.txt
202    readStart = getCurrentTime();
203    numRead = read(fileDescriptor, &buffer, 10000);
204    readEnd = getCurrentTime();
205    printf("Time to execute read() for %d bytes: %ld\n", numRead, (readEnd -
        readStart));
206
207    // Read 100000 bytes from read.txt
208    readStart = getCurrentTime();
209    numRead = read(fileDescriptor, &buffer, 100000);
210    readEnd = getCurrentTime();
211    printf("Time to execute read() for %d bytes: %ld\n", numRead, (readEnd -
        readStart));
212
213    // Preparation for mkdir() measurements
214    unsigned long mkdirStart;
215    unsigned long mkdirEnd;
216    int success;
217
218    // Make a directory with all access permissions given
219    mkdirStart = getCurrentTime();
220    success = mkdir("allAccessPermissionsGiven", ACCESSPERMS);
221    mkdirEnd = getCurrentTime();
222    printf("\nTime to execute mkdir() giving all access permissions: %ld\n",
        mkdirEnd - mkdirStart);
223    printf("Successful creation (0 is OK if -1 the directory probably exists): %d\
        n", success);
224
225    // Make a directory with only read permissions given
226    mkdirStart = getCurrentTime();
227    success = mkdir("onlyReadPermissionsGiven", S_IRUSR | S_IRGRP | S_IROTH);
228    mkdirEnd = getCurrentTime();
229    printf("Time to execute mkdir() giving only read permissions: %ld\n", mkdirEnd
        - mkdirStart);
230    printf("Successful creation (0 is OK if -1 the directory probably exists): %d\
        n", success);
231
232    return 0;
233 }
234
235 /**
236  * Returns the current time in microseconds. The current
```

```
237 * time corresponds to the time elapsed from the starting
238 * point used by the gettimeofday() function.
239 * @return currentTime: The current time in microseconds
240 */
241 unsigned long getCurrentTime() {
242     struct timeval timeValue;
243     gettimeofday(&timeValue, NULL);
244     unsigned long currentTime = timeValue.tv_usec; // microseconds part of the
        struct
245     currentTime += timeValue.tv_sec * 1e6; // add the seconds part of the struct
246     return currentTime;
247 }
```

The command line output:

```
249 efe@efe-VirtualBox:~/Desktop$ make
250 gcc -Wall -g -o cost cost.c
251 efe@efe-VirtualBox:~/Desktop$ ./cost
252
253 Time measurements for different system calls in microseconds follows:
254
255 Time to execute getpid(): 2
256 Process ID is: 4160
257
258 Time to execute open() to create a new .txt file is: 33
259 The value of fileDescriptor is 3:
260
261 Time to execute write() for 100 bytes: 13
262 Time to execute write() for 1000 bytes: 3
263 Time to execute write() for 10000 bytes: 7
264 Time to execute write() for 100000 bytes: 42
265
266 Time to execute read() for 100 bytes: 1
267 Time to execute read() for 1000 bytes: 1
268 Time to execute read() for 10000 bytes: 4
269 Time to execute read() for 100000 bytes: 31
270
271 Time to execute mkdir() giving all access permissions: 25
272 Successful creation (0 is OK if -1 the directory probably exists): 0
273 Time to execute mkdir() giving only read permissions: 14
274 Successful creation (0 is OK if -1 the directory probably exists): 0
```

For the make command, The Makefile provided in the homework assignment sheet is used as it is.